

EXP NO: 7 — VERSION 1

GENERATIVE MODELS WITH GANS: CREATING AND TRAINING A GENERATIVE ADVERSARIAL NETWORK

AIM

To construct and train a Generative Adversarial Network (GAN) using TensorFlow/Keras on the MNIST dataset to generate new synthetic images of handwritten digits.

PROCEDURE

1. Import the required TensorFlow, Keras, NumPy and Matplotlib libraries.
2. Load the MNIST dataset and normalize the images to the range -1 to 1 .
3. Define the Generator network using dense, batch normalization and transposed convolution layers.
4. Define the Discriminator network using convolution, dropout and dense layers.
5. Define the generator and discriminator loss functions.
6. Create Adam optimizers for both networks.
7. Train the GAN using real MNIST images and generated images.
8. Generate and display synthetic digit images at different training stages.

CODE

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
import os

(x_train, _), (_, _) = mnist.load_data()

x_train = x_train.reshape(-1, 28, 28, 1).astype("float32")
```

```
x_train = (x_train - 127.5) / 127.5
```

```
latent_dim = 100
```

```
batch_size = 256
```

```
epochs = 30
```

```
num_examples = 16
```

```
generator = keras.Sequential([  
    layers.Input(shape=(latent_dim,)),  
    layers.Dense(7 * 7 * 256, use_bias=False),  
    layers.BatchNormalization(),  
    layers.LeakyReLU(),  
    layers.Reshape((7, 7, 256)),  
    layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding="same", use_bias=False),  
    layers.BatchNormalization(),  
    layers.LeakyReLU(),  
    layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding="same", use_bias=False),  
    layers.BatchNormalization(),  
    layers.LeakyReLU(),  
    layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding="same", activation="tanh")  
])
```

```
discriminator = keras.Sequential([  
    layers.Input(shape=(28, 28, 1)),  
    layers.Conv2D(64, (5, 5), strides=(2, 2), padding="same"),  
    layers.LeakyReLU(),  
    layers.Dropout(0.3),
```

```

layers.Conv2D(128, (5, 5), strides=(2, 2), padding="same"),
layers.LeakyReLU(),
layers.Dropout(0.3),
layers.Flatten(),
layers.Dense(1, activation="sigmoid")
])

loss_function = keras.losses.BinaryCrossentropy()

def generator_loss(fake_output):
    return loss_function(tf.ones_like(fake_output), fake_output)

def discriminator_loss(real_output, fake_output):
    real_loss = loss_function(tf.ones_like(real_output), real_output)
    fake_loss = loss_function(tf.zeros_like(fake_output), fake_output)
    return real_loss + fake_loss

generator_optimizer = keras.optimizers.Adam(learning_rate=0.0001)
discriminator_optimizer = keras.optimizers.Adam(learning_rate=0.0001)

@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, latent_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

```

```

real_output = discriminator(images, training=True)
fake_output = discriminator(generated_images, training=True)

gen_loss = generator_loss(fake_output)
disc_loss = discriminator_loss(real_output, fake_output)

gen_gradients = gen_tape.gradient(
    gen_loss,
    generator.trainable_variables
)

disc_gradients = disc_tape.gradient(
    disc_loss,
    discriminator.trainable_variables
)

generator_optimizer.apply_gradients(
    zip(gen_gradients, generator.trainable_variables)
)

discriminator_optimizer.apply_gradients(
    zip(disc_gradients, discriminator.trainable_variables)
)

return gen_loss, disc_loss

dataset = tf.data.Dataset.from_tensor_slices(x_train)

```

```
dataset = dataset.shuffle(60000).batch(batch_size, drop_remainder=True)
```

```
seed = tf.random.normal([num_examples, latent_dim])
```

```
def generate_images(epoch):
```

```
    predictions = generator(seed, training=False)
```

```
    predictions = (predictions + 1) / 2
```

```
    plt.figure(figsize=(6, 6))
```

```
    for i in range(num_examples):
```

```
        plt.subplot(4, 4, i + 1)
```

```
        plt.imshow(predictions[i, :, :, 0], cmap="gray")
```

```
        plt.axis("off")
```

```
    plt.suptitle("Generated Images - Epoch " + str(epoch))
```

```
    plt.tight_layout()
```

```
    plt.show()
```

```
for epoch in range(epochs):
```

```
    generator_losses = []
```

```
    discriminator_losses = []
```

```
    for image_batch in dataset:
```

```
        gen_loss, disc_loss = train_step(image_batch)
```

```
        generator_losses.append(gen_loss.numpy())
```

```
        discriminator_losses.append(disc_loss.numpy())
```

```

print(
    "Epoch:",
    epoch + 1,
    "Generator Loss:",
    round(np.mean(generator_losses), 4),
    "Discriminator Loss:",
    round(np.mean(discriminator_losses), 4)
)

if (epoch + 1) % 5 == 0:
    generate_images(epoch + 1)

```

```
generate_images(epochs)
```


OUTPUT

```

Epoch: 1 Generator Loss: 0.6491 Discriminator Loss: 1.1856
Epoch: 2 Generator Loss: 0.8037 Discriminator Loss: 1.2403
Epoch: 3 Generator Loss: 0.8976 Discriminator Loss: 1.2006
Epoch: 4 Generator Loss: 0.7084 Discriminator Loss: 1.3152
Epoch: 5 Generator Loss: 0.8998 Discriminator Loss: 1.2251

```

Generated Images - Epoch 5




```

Epoch: 6 Generator Loss: 0.8532 Discriminator Loss: 1.2955
Epoch: 7 Generator Loss: 0.8666 Discriminator Loss: 1.2559
Epoch: 8 Generator Loss: 0.9358 Discriminator Loss: 1.2021
Epoch: 9 Generator Loss: 0.9055 Discriminator Loss: 1.2172
Epoch: 10 Generator Loss: 0.8858 Discriminator Loss: 1.2776


```

Generated Images - Epoch 10



```
Epoch: 11 Generator Loss: 0.8721 Discriminator Loss: 1.271
Epoch: 12 Generator Loss: 0.9429 Discriminator Loss: 1.2113
Epoch: 13 Generator Loss: 1.0059 Discriminator Loss: 1.1668
Epoch: 14 Generator Loss: 0.9562 Discriminator Loss: 1.2058
Epoch: 15 Generator Loss: 1.0847 Discriminator Loss: 1.1247
```

Generated Images - Epoch 15



```
Epoch: 16 Generator Loss: 1.0448 Discriminator Loss: 1.1707
Epoch: 17 Generator Loss: 1.0384 Discriminator Loss: 1.1555
Epoch: 18 Generator Loss: 1.0241 Discriminator Loss: 1.1636
Epoch: 19 Generator Loss: 1.0041 Discriminator Loss: 1.0099
```

RESULT

Thus, the Generative Adversarial Network was successfully constructed and trained using the MNIST dataset. The Generator learned to create synthetic handwritten digit images, while the Discriminator learned to distinguish between real and generated images. The generated images became more realistic as training progressed.